

Lecture 11

computer clusters



Course: Practical Bioinformatics (BIOL 4220)
Instructor: Michael Landis
Email: michael.landis@wustl.edu



Lecture 11 outline

Last time: regex substitutions

This time: clusters

cluster basics

- cluster anatomy
- using clusters
- sharing data w/ clusters

Servers and clusters

A ***server*** is a computer with ample resources that are shared among many ***clients***

A ***cluster*** is a group of servers that are configured to distribute tasks that involve large numbers of jobs and/or parallelized jobs

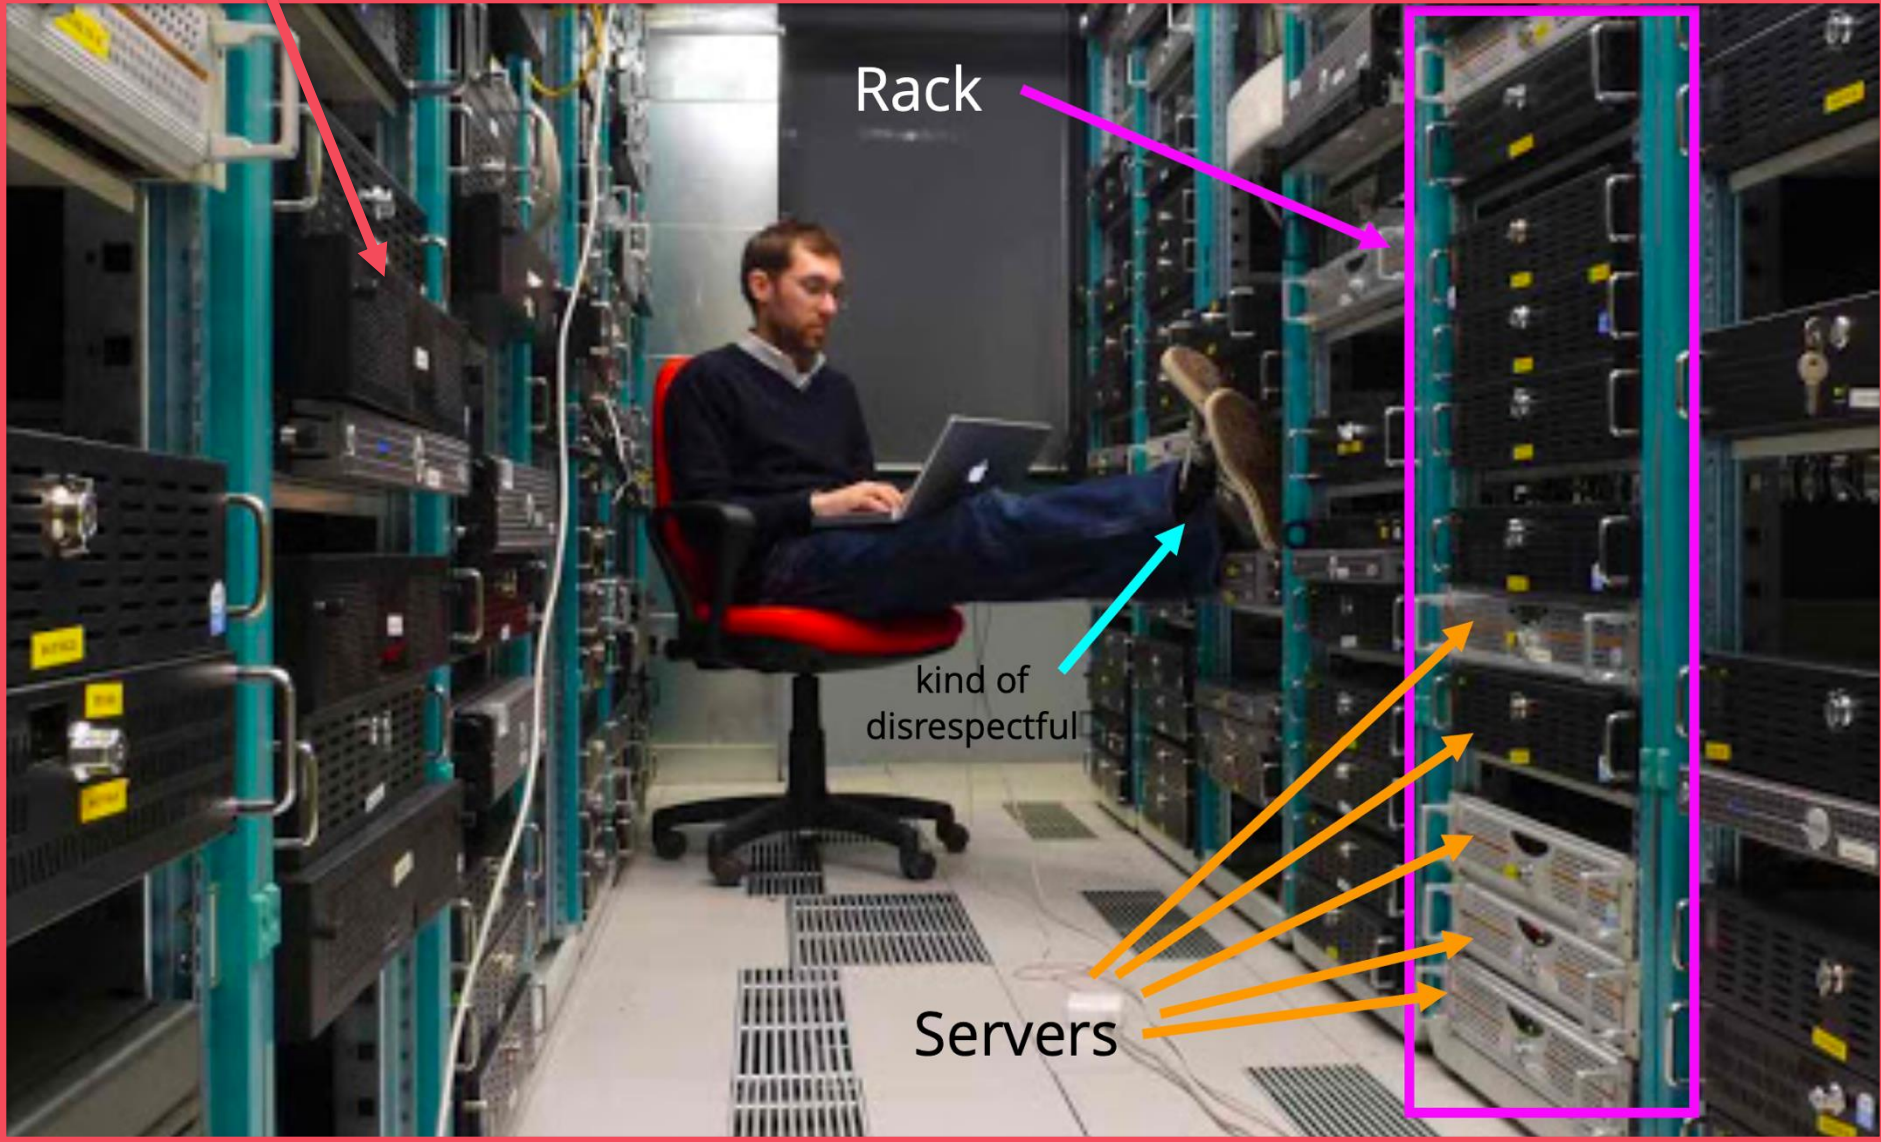
Information Technology (IT) departments manage security, access, maintenance, and expansion of servers and clusters

Cluster

Rack

kind of
disrespectful

Servers



Why clusters?

Large institutions that rely on computation invest in clusters for many reasons:

- massive parallelization of jobs
- large shared storage
- more efficient (less idling)
- economy of scale
- greater hardware capacity
- centralized security

Cluster anatomy

Common cluster features:


- Many ***users***, each with individual needs
- ***Private network*** to secure resources and monitor bandwidth
- ***Nodes*** are servers with well-defined roles, depending on hardware and network access
 - users log-in to and submit jobs through the ***head node***
 - the ***scheduler*** determines when to process each ***job***
 - jobs are processed by the large pool of ***compute nodes***
 - each compute node generally has 16-64 ***compute cores***
 - cores are the units of ***parallelization***
 - some jobs might request nodes with special hardware, such as ***graphics processing unit cards (GPUs)***

Cluster anatomy

Common cluster features (cont'd):

- All users and nodes share access to ***network storage***
 - ***OS storage*** for e.g. utilities, libraries, home directories; small, fast, permanent, and free
 - ***Scratch storage*** for job output, *etc.* that is periodically deleted; large, slow, temporary, and free
 - ***Persistent storage*** for large irreplaceable datasets; large, slow, permanent, and costly
- Nodes, storage devices, network devices, *etc.* are installed in ***racks*** in ***data centers***, where IT manages network, backup, storage, power, *etc.*

Cluster anatomy



Five fastest clusters

Rank	System	Cores	Rmax (PFlop/s)	Rpeak (PFlop/s)	Power (kW)
1	Frontier - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE DOE/SC/Oak Ridge National Laboratory United States	8,699,904	1,206.00	1,714.81	22,786
2	Aurora - HPE Cray EX - Intel Exascale Compute Blade, Xeon CPU Max 9470 52C 2.4GHz, Intel Data Center GPU Max, Slingshot-11, Intel DOE/SC/Argonne National Laboratory United States	9,264,128	1,012.00	1,980.01	38,698
3	Eagle - Microsoft NDv5, Xeon Platinum 8480C 48C 2GHz, NVIDIA H100, NVIDIA Infiniband NDR, Microsoft Azure Microsoft Azure United States	2,073,600	561.20	846.84	
4	Supercomputer Fugaku - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu RIKEN Center for Computational Science Japan	7,630,848	442.01	537.21	29,899
5	LUMI - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE EuroHPC/CSC Finland	2,752,704	379.70	531.51	7,107

PFlop/s, 1e15 floating point operations/s

<https://www.top500.org/lists/top500/2024/06/>

WashU RIS cluster

- Base system
 - 5,000 Intel Cascade Lake Cores
 - 120 NVIDIA Tesla V100 GPUs
 - 300TB DDN high-performance scratch space
 - 100Gbit Mellanox HDR Network

Rough estimate?

CPU ~0.57 PFlop/s

GPU ~1.88 PFlop/s

WashU cluster resources

- Resources webpage
<https://ris.wustl.edu/resources/>
- Manual
<https://docs.ris.wustl.edu/index.html>
- Knowledgebase
<https://washu.atlassian.net/jira/servicedesk/projects/ITSD/knowledge/articles>
- Help desk
<http://servicedesk.ris.wustl.edu/>

Using clusters

Typical workflow for cluster

1. Connect to cluster network (e.g. VPN)
2. Log into cluster through *head node* using *ssh*
3. Set up files and profile on cluster
4. Configure scripts on cluster to work with programs and filesystem
5. Submit jobs to *scheduler*
6. Wait for *compute nodes* to run jobs
7. Retrieve output from cluster

We'll practice this workflow in lab

Job scheduler

Clusters are endowed with an extensive array of hardware, *e.g.* CPU, GPU, memory, storage

The ***job scheduler*** ensure that resources are shared fairly among users by managing multiple ***job queues***

Users submit their job(s) to a queue, the queue then processes jobs based on ***queue policies***



tragedy of the commons

Job scheduler

Example scenario

Both jobs require 6400 CPU-hours

- User 1 submits 6400x 1-core 1-hour jobs
- User 2 submits 10x 64-core 10-hour jobs

Easier to secure resources for User 1 (many, small) than User 2 (few, large)

Scheduler priorities, managed through policy:

- User 1 jobs decrease in priority as more jobs are run
- User 2 jobs increase in priority the longer they wait

Job scheduler

Different clusters use different software for job scheduling.

Examples: Slurm, Torque, LSF

How users interact with the job scheduler varies in terms of syntax, but workflow is fairly consistent across platforms:

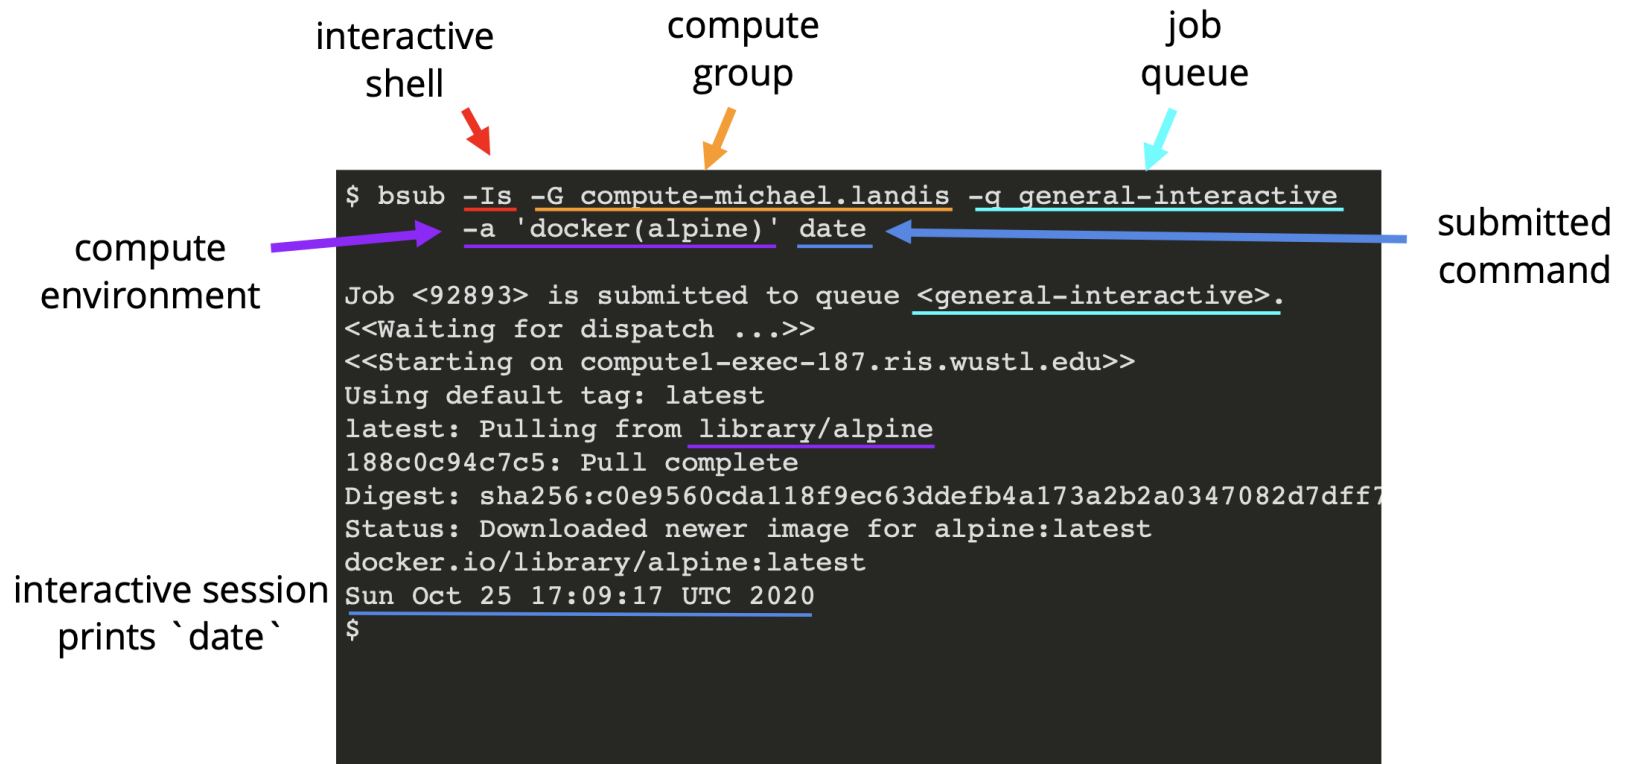
- submit jobs
- monitor jobs
- cancel jobs

The RIS cluster uses LSF:

<https://www.ibm.com/docs/en/spectrum-lsf/10.1.0>

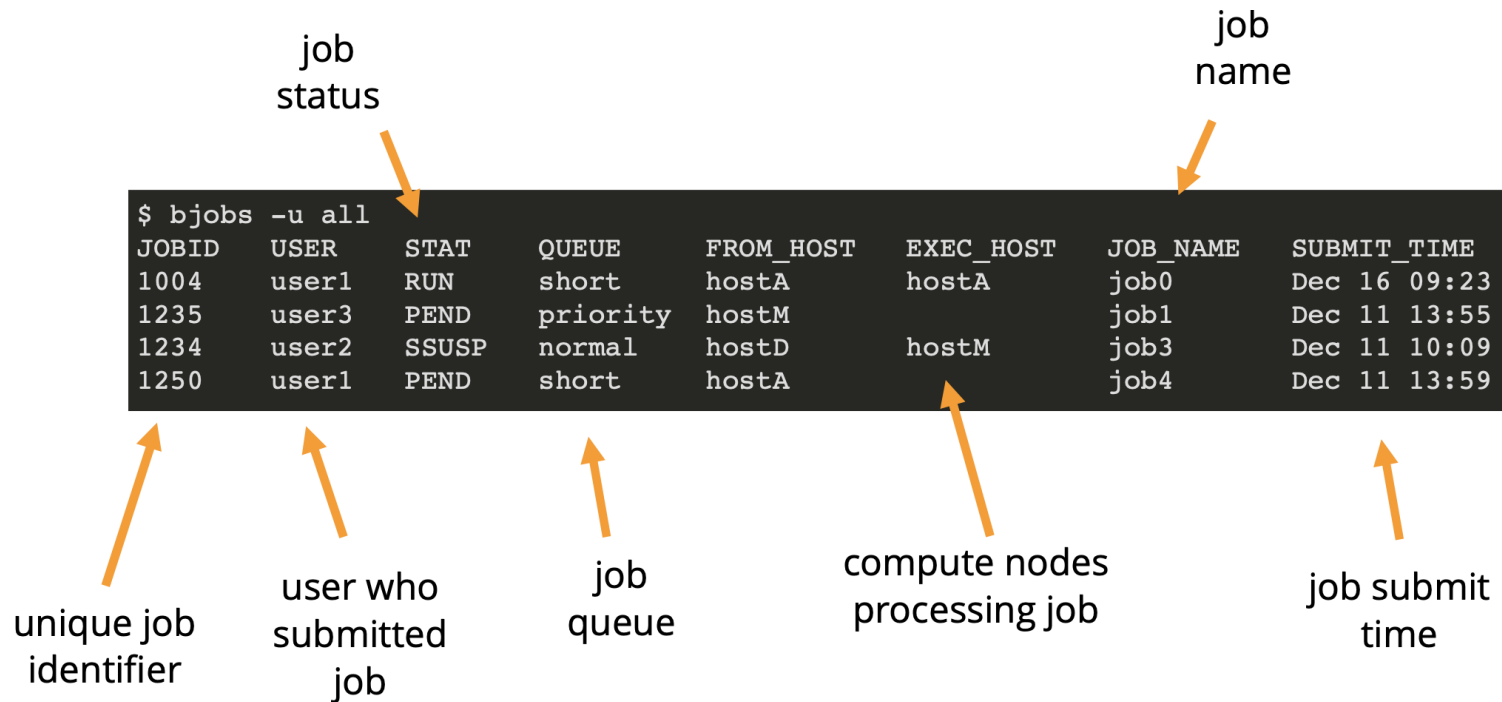
bsub

Submit a job to queue



bjobs

List all queued jobs



The diagram illustrates the output of the `bjobs` command. It features a dark rectangular box containing a terminal-style output. Orange arrows point from descriptive labels to specific columns in the table. The labels are: 'job status' (points to STAT), 'job name' (points to JOB_NAME), 'unique job identifier' (points to JOBID), 'user who submitted job' (points to USER), 'job queue' (points to QUEUE), 'compute nodes processing job' (points to EXEC_HOST), and 'job submit time' (points to SUBMIT_TIME).

JOBID	USER	STAT	QUEUE	FROM_HOST	EXEC_HOST	JOB_NAME	SUBMIT_TIME
1004	user1	RUN	short	hostA	hostA	job0	Dec 16 09:23
1235	user3	PEND	priority	hostM		job1	Dec 11 13:55
1234	user2	SSUSP	normal	hostD	hostM	job3	Dec 11 10:09
1250	user1	PEND	short	hostA		job4	Dec 11 13:59

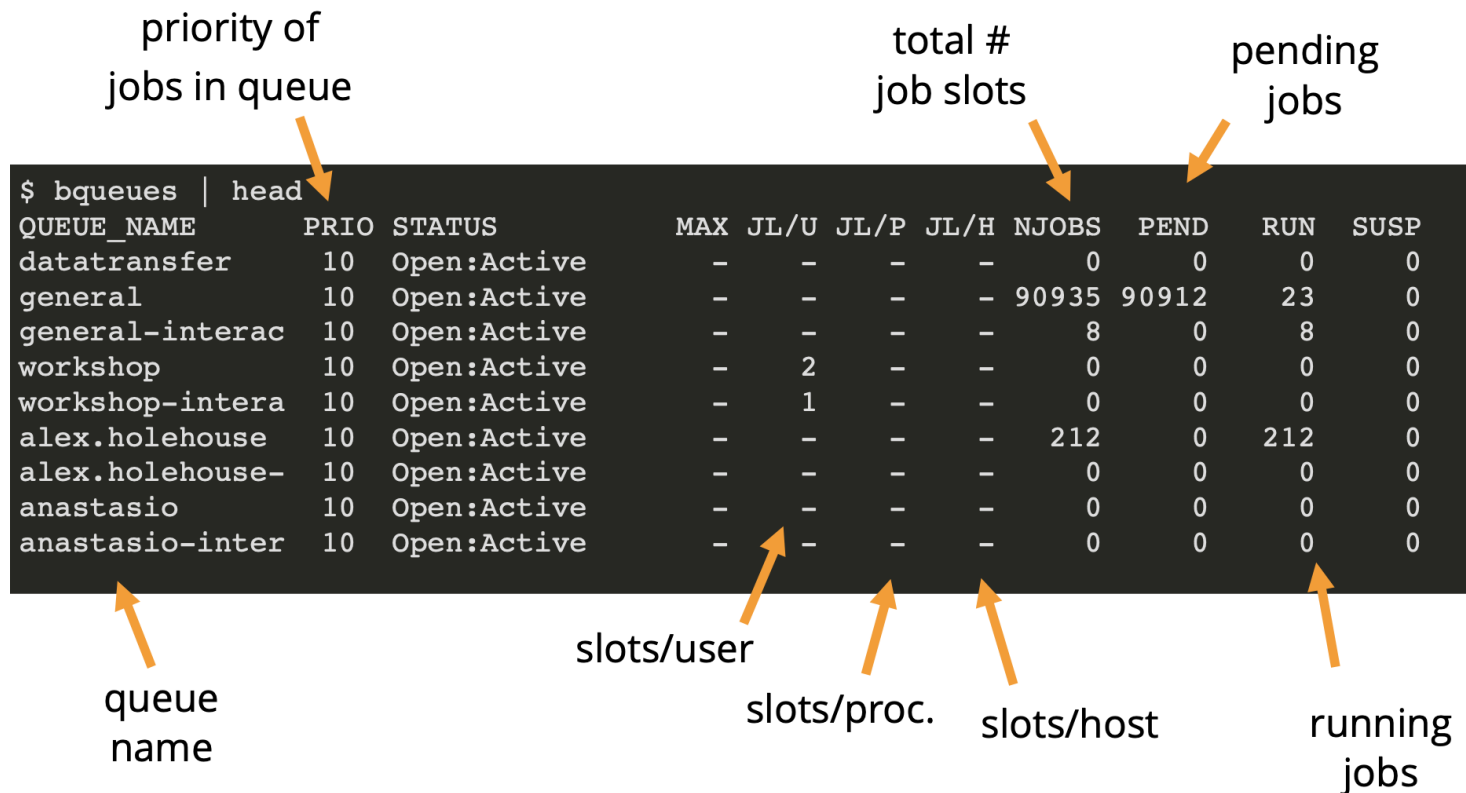
bkill

Terminate a job in the queue

```
$ bsub -G compute-michael.landis \  
> -q general \  
> -a 'docker(alpine)' date  
→ Job <92898> is submitted to queue <general>.  
# list jobs  
$ bjobs  
→  
JOBID    USER      STAT  QUEUE      FROM_HOST   EXEC_HOST   JOB_NAME   SUBMIT_TIME  
92898    michael   PEND   general     computel-cl  date        Oct 25 12:23  
# kill job  
$ bkill 92898  
→ Job <92898> is being terminated  
# job no longer running  
$ bjobs  
No unfinished job found
```

bqueues

List all queues for cluster
Helps identify underused queues



priority of jobs in queue

total # job slots

pending jobs

```
$ bqueues | head
```

QUEUE_NAME	PRIO	STATUS	MAX	JL/U	JL/P	JL/H	NJOBS	PEND	RUN	SUSP
datatransfer	10	Open:Active	-	-	-	-	0	0	0	0
general	10	Open:Active	-	-	-	-	90935	90912	23	0
general-interac	10	Open:Active	-	-	-	-	8	0	8	0
workshop	10	Open:Active	-	2	-	-	0	0	0	0
workshop-intera	10	Open:Active	-	1	-	-	0	0	0	0
alex.holehouse	10	Open:Active	-	-	-	-	212	0	212	0
alex.holehouse-	10	Open:Active	-	-	-	-	0	0	0	0
anastasio	10	Open:Active	-	-	-	-	0	0	0	0
anastasio-inter	10	Open:Active	-	-	-	-	0	0	0	0

queue name

slots/user

slots/proc.

slots/host

running jobs

scp

Secure copy allows you to transfer files between local and remote filesystems

```
# copy file1.txt FROM remote server TO local machine
$ scp mlandis@some.server.org:/home/mlandis/file.txt .
mlandis@some.server.org's password: <password>
file1.txt                                100%  413    22.7KB/s   00:00

# copy file2.txt TO remote server FROM local machine
$ scp file.txt mlandis@some.server.org:/home/mlandis
mlandis@some.server.org's password: <password>
file2.txt                                100%  777    22.9KB/s   00:00

# recursively copy my_dir TO remote server FROM local machine
$ scp -r my_dir mlandis@some.server.org:/home/mlandis

michael.landis@computel-client-1.ris.wustl.edu's password:
file1.txt                                100%    0     0.0KB/s   00:00
file2.txt                                100%  413    24.0KB/s   00:00

# remote server now contains the files/directories:
# /home/mlandis/my_dir/file1.txt
# /home/mlandis/my_dir/file2.txt
```

Overview for Lab 11