

Lecture 21

SciPy



Course: Practical Bioinformatics (BIOL 4220)
Instructor: Michael Landis
Email: michael.landis@wustl.edu



Lecture 21 outline

Last time: NumPy

This time: SciPy

- overview of SciPy
- survey of SciPy modules



SciPy is an open source ecosystem that extends Python's functionality for math, science, and engineering

This ecosystem includes its open source software, its community, and its conferences

The SciPy library itself provides methods for signal processing, optimization, integration, statistics, and more

SciPy library organization

→	<u>cluster</u>	Clustering algorithms
→	<u>constants</u>	Physical and mathematical constants
	<u>fftpack</u>	Fast Fourier Transform routines
→	<u>integrate</u>	Integration and ordinary differential equation solvers
	<u>interpolate</u>	Interpolation and smoothing splines
	<u>io</u>	Input and Output
	<u>linalg</u>	Linear algebra
	<u>ndimage</u>	N-dimensional image processing
	<u>odr</u>	Orthogonal distance regression
→	<u>optimize</u>	Optimization and root-finding routines
	<u>signal</u>	Signal processing
	<u>sparse</u>	Sparse matrices and associated routines
→	<u>spatial</u>	Spatial data structures and algorithms
	<u>special</u>	Special functions
→	<u>stats</u>	Statistical distributions and functions

scipy.constants

```
>>> from scipy import constants
>>> constants.pi                # pi
3.141592653589793
>>> constants.golden            # golden ratio, (1+50.5)/2
1.618033988749895
>>> constants.Avogadro          # Avogadro's number
6.022140857e+23
>>> constants.speed_of_light    # speed of light in vacuum
299792458.0
>>> constants.electron_mass     # mass of electron
9.10938356e-31
>>> constants.proton_mass       # mass of proton
1.672621898e-27
>>> constants.neutron_mass      # mass of neutron
1.674927471e-27
>>> scipy.constants.physical_constants # returns dict of (values, units, precision)
{'Wien displacement law constant': (0.0028977685, 'm K', 5.1e-09),
 'atomic unit of 1st hyperpolarizability': (3.20636151e-53, 'C3 m3 J-2', 2.8e-60),
 'atomic unit of 2nd hyperpolarizability': (6.2353808e-65, 'C4 m4 J-3', 1.1e-71),
 'atomic unit of electric dipole moment': (8.47835309e-30, 'C m', 7.3e-37),
 'atomic unit of electric polarizability': (1.648777274e-41, 'C2 m2 J-1', 1.6e-49),
 'atomic unit of electric quadrupole moment': (4.48655124e-40, 'C m2', 3.9e-47),
 ...
```

*Mathematical and (mostly) physical
constants, units, and precisions*

scipy.stats

Provides methods for numerous probability distributions and a wide variety of statistical operations

- probability distributions
- statistical tests
- frequency statistics
- summary statistics
- statistical distances
- contingency tables
- kernel density estimators

distributions in *scipy.stats*

Probability distributions assign relative frequencies to possible outcomes for random experiment – *e.g. a coin flip*

All *scipy.stats* distributions offer the same set of base functions:

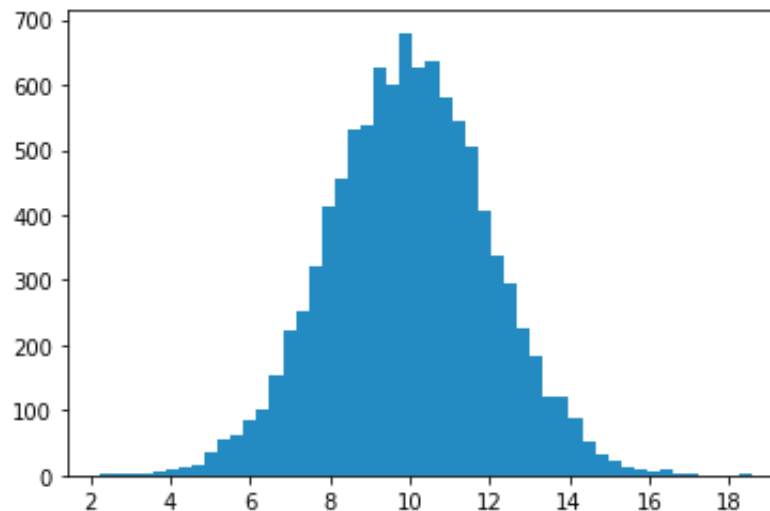
- simulate random variates
- compute distribution probabilities for data
- compute moments (mean, variance, etc.)
- fit distribution parameters to data

distributions in *scipy.stats*

```
>>> from scipy import stats
>>> # distribution object
>>> stats.norm
<scipy.stats._continuous_distns.norm_gen object at 0x7fdf98b67d68>
>>> # generate 4 normal RVs
>>> x = stats.norm.rvs(loc=10, scale=2, size=4)
array([10.11179033, 10.10902411, 10.65111753, 10.32368948])
>>> stats.norm.pdf(x=x, loc=10, scale=2)
array([0.19915978, 0.19917499, 0.18917553, 0.19687573])
>>> # return mean, variance, skewness, kurtosis (mvsk)
>>> stats.norm.stats(loc=10, scale=2, moments='mvsk')
(array(10.), array(4.), array(0.), array(0.))
>>> # generate 1000 normal RVs with mean=10, scale=2
>>> y = stats.norm.rvs(loc=10, scale=2, size=1000)
>>> # estimate mean and scale from simulated data
>>> stats.norm.fit(data=y, loc=50, scale=9)
(9.939835173664239, 2.004585580838588)
```


distributions in *scipy.stats*

```
>>> from scipy import stats
>>> import matplotlib
>>> import matplotlib.pyplot as plt
>>> # simulate 10000 RVs, mean=10, scale=2
>>> y = stats.norm.rvs(loc=10,scale=2,size=10000)
>>> y
array([13.23832222, 13.05731999, 7.30259037, ..., 9.8595785,
        9.46715211, 9.98579946])
>>> fig,ax = plt.subplots()
>>> p = ax.hist(y,bins=50)
>>> fig.show()
```



distributions in *scipy.stats*

Roughly 100 distributions available,
each with different properties

Type	Name	Class	Use
Continuous	uniform	<i>stats.uniform</i>	"flat" over interval
	normal	<i>stats.norm</i>	random sums
	exponential	<i>stats.expon</i>	events with rates
	beta	<i>stats.beta</i>	flexible on [0,1]
Discrete	bernoulli	<i>stats.bernoulli</i>	single coin-flip
	binomial	<i>stats.binom</i>	many coin-flips
	Poisson	<i>stats.poisson</i>	# events w/ rates

summary statistics in *scipy.stats*

Use ***summary statistics*** to describe simple properties of a data sample, e.g. *the sample mean*

```
>>> from scipy import stats
>>> import numpy as np
>>> x = stats.norm.rvs(loc=10, scale=2, size=1000)
>>> np.mean(x)                # sample mean
9.907750191507521
>>> stats.hmean(x)            # harmonic mean
9.4599529015089185
>>> stats.gmean(x)            # geometric mean
9.6924828402578722
>>> stats.sem(x)              # standard error of sample mean
0.06353084579162438
>>> stats.skew(x)             # sample skew (asymmetry)
0.011612124738411802
>>> stats.kurtosis(x)         # sample kurtosis (fat-tailedness)
-0.0055425614273967305
>>> stats.describe(x)         # give summary of sample data
DescribeResult(nobs=1000,
               minmax=(3.7313631618255805, 16.221776607589973),
               mean=9.907750191507521,
               variance=4.0361683669991582,
               skewness=0.011612124738411802,
               kurtosis=-0.0055425614273967305)
```

statistical tests in *scipy.stats*

Many statistical tests that are used to analyze biological data are available, including:

- t-tests: *ttest_1samp*, *ttest_ind*, *ttest_rel*
 - Chi-square: *chisquare*
 - Kolmogorov-Smirnov: *kstest*, *ks_1samp*, *ks_2samp*
 - Mann-Whitney: *mannwhitneyu*
 - Cressie-Read power divergence: *power_divergence*
 - Wilcoxon signed-rank test: *wilcoxon*
 - Kruskal-Wallis H-test: *kruskal*
- ...and dozens more*

Linear regression in *scipy.stats*


Fit data to line using slope and intercept parameters

```
>>> # import numpy and scipy
>>> import numpy as np
>>> import scipy as sp
>>> # simulate dataset
>>> np.random.seed(seed=12345)
>>> x = sp.stats.uniform.rvs(size=100, loc=0, scale=1)
>>> y = 1.6*x + sp.stats.uniform.rvs(size=100, loc=0, scale=1)
>>> # get results from linear regression,  $y \sim x$ 
>>> res = sp.stats.linregress(x, y)
>>> res
LinregressResult(slope=1.4431933040436178, intercept=0.5860274273931283,
rvalue=0.8088733612574854, pvalue=2.4706574634065375e-24, stderr=0.1059730971193916,
intercept_stderr=0.06594536829897774)
>>> print(f"R-squared: {res.rvalue**2:.6f}")
R-squared: 0.65427
```

regression

```
>>> # plot results
>>> import matplotlib.pyplot as plt
>>> plt.plot(x, y, 'o', label='original data')
>>> plt.plot(x, res.intercept + res.slope*x, 'r',
label='fitted line')
>>> plt.legend()
>>> plt.show()
```

plotting




KDEs in *scipy.stats*

Kernel density estimators (KDEs) approximate sample data as a continuous probability density

```
>>> from scipy import stats
>>> import matplotlib
>>> import matplotlib.pyplot as plt
>>> a = stats.norm.rvs(loc=12,scale=2,size=10000) # create samples, a
>>> b = stats.norm.rvs(loc=6,scale=1,size=10000) # create samples, b
>>> z = np.append(a, b) # mix samples a and b
>>> f = stats.gaussian_kde(z) # fit KDE object to data, z
>>> x = np.arange(0,20,0.1) # range of input values for f()
>>> f(x) # density for each value, f(pos)
array([1.33471717e-11, 4.45526152e-11, 1.42288247e-10, 4.34870069e-10,
        1.27215449e-09, 3.56303164e-09, 9.55708453e-09, 2.45585700e-08,
        6.04813171e-08, 1.42815299e-07, 3.23510616e-07, 7.03432155e-07,
        ...
>>> fig,ax = plt.subplots() # create empty plot
>>> h = ax.hist(z,bins=50,density=True) # plot histogram (normalized)
>>> ax.plot(x, f(x)) # plot KDE curve
>>> plt.show() # show plot
```

KDEs in *scipy.stats*


Kernel density estimators (KDEs) approximate sample data as a continuous probability density



scipy.optimize

```
>>> import numpy as np
>>> from scipy.optimize import minimize
>>> def rosen(x):
...     # Rosenbrock function, for a=1 and b=100
...     # f(x,y) = (1-x)^2 + 100*(y-x^2)^2
...     return (1.0-x[0])**2 + 100.0*(x[1]-x[0]**2.0)**2.0
>>>
>>> rosen( [1.0, 1.0] )           # optimal value, f(1,1)=0
0.0
>>> rosen( [1.01, 1.01] )        # worse value
0.010300999999999994
>>> rosen( [1.05, 1.05] )        # worse value
0.27812499999999999
>>> x0 = np.array([1.3, 0.7]) # initial guess
>>> # find the optimal value for x
>>> res = minimize(rosen, x0, method='nelder-mead',
...                 options={'xatol': 1e-8, 'disp': True})
Optimization terminated successfully.
Current function value: 0.000000
Iterations: 79
Function evaluations: 150
>>> print(res.x)                  # estimate for x
[1. 1.]
>>> print(res.fun)                # minimum, approx 0.0
3.3736077629532093e-18
```

$$f(x, y) = (a - x)^2 + b(y - x^2)^2$$
$$a = 1, b = 100$$
$$f(1, 1) = 0$$



scipy.integrate

Various methods to integrate functions,
fixed samples and ODEs


```
>>> from scipy import integrate
>>> # define function to integrate
>>> def x2(x):
...     return x**2
...
>>> # numerical integration using quadrature
>>> y, err = integrate.quad(x2, 0, 4)
>>> y
21.333333333333336      # numerical integrand
>>> err
2.368475785867001e-13  # numerical error
>>>
>>> print(4**3 / 3.)    # analytical result
21.3333333333
```

$$\begin{aligned} I &= \int_0^4 x^2 dx \\ &= F(4) - F(0) \\ &= \frac{4^3}{3} - \frac{0^3}{3} = 21\frac{1}{3} \end{aligned}$$

scipy.clustering

Assorted functions for inferring latent structures in data, e.g. *k-means*, *vector-quantization*, *hierarchical clustering*

```
>>> import numpy as np
>>> from scipy.cluster.vq import vq, kmeans, whiten
>>> import matplotlib.pyplot as plt
>>> # 50 pts at (0,0)
>>> a = np.random.multivariate_normal([ 0, 0],
... [[ 4, 1], [1, 4]],
... size=50)
...
>>> # 50 pts at (30,10)
>>> b = np.random.multivariate_normal([30, 10],
... [[10, 2], [2, 1]],
... size=50)
...
>>> features = np.concatenate((a, b)) # x = [ a, b ]
>>> x = whiten(features) # normalize data
>>> y, e = kmeans(x, 2) # infer k=2 means
>>> plt.scatter(x[:, 0], x[:, 1]) # plot data, x
>>> plt.scatter(y[:, 0], y[:, 1], # plot means, y
...             c='r')
>>> plt.show()
```




scipy.spatial.ConvexHull

polygons to quantify the spread/shape of data

```
>>> # import numpy/scipy
>>> import numpy as np
>>> import scipy as sp
>>> # simulate dataset
>>> np.random.seed(seed=12345)
>>> points = sp.stats.uniform.rvs(size=(30,2))
>>> # find convex hull of points
>>> hull = sp.spatial.ConvexHull(points)
>>> hull.area
3.008424594700238
>>> hull.points[hull.vertices,:]
```

Vertex Index	x	y
0	0.846943	0.37809954
1	0.9174485	0.80358537
2	0.62491062	0.89256119
3	0.19608178	0.9354918
4	0.10410446	0.71212017
5	0.03714544	0.03413158
6	0.37113194	0.04055327
7	0.4249371	0.0743783



```
>>> # import numpy/scipy
>>> import matplotlib.pyplot as plt
>>> # plot raw data
>>> plt.plot(points[:,0], points[:,1], 'o')
>>> # plot convex hull
>>> for simplex in hull.simplices:
+++     plt.plot(points[simplex, 0], points[simplex, 1], 'k-')
>>> plt.plot(points[hull.vertices,0], points[hull.vertices,1], 'r--', lw=2)
>>> plt.plot(points[hull.vertices[0],0], points[hull.vertices[0],1], 'ro')
>>> plt.show()
```

Overview for Lab 21